

# FlowWalker: A Memory-efficient and High-performance GPU-based Dynamic Graph Random Walk Framework

🎤 **Junyi Mei**<sup>1</sup>, Shixuan Sun<sup>1</sup>, Chao Li<sup>1</sup>, Cheng Xu<sup>1</sup>, Cheng Chen<sup>2</sup>, Yibo Liu<sup>1</sup>, Jing Wang<sup>1</sup>,  
Cheng Zhao<sup>2</sup>, Xiaofeng Hou<sup>1</sup>, Minyi Guo<sup>1</sup>, Bingsheng He<sup>3</sup>, Xiaoliang Cong<sup>2</sup>

<sup>1</sup>Shanghai Jiao Tong University

<sup>2</sup>ByteDance Inc

<sup>3</sup>National University of Singapore




SHANGHAI JIAO TONG  
UNIVERSITY



# Importance of Graph Random Walk (RW)


RW has been used in many real-world applications:

- Social network analysis
- Recommendation system
- Knowledge graph
- ...




## Graph neural network

- ❑ GraphSaint
- ❑ PinSage
- ❑ ...



## Graph embedding

- ❑ DeepWalk
- ❑ Node2vec
- ❑ ...



## Graph ranking

- ❑ Personalized PageRank
- ❑ SimRank
- ❑ ...

Figure source:

[1]. William L. Hamilton, Rex Ying, and Jure Leskovec. 2017. Inductive representation learning on large graphs. In Proceedings of the 31st International Conference on Neural Information Processing Systems (NIPS'17).

[2]. [https://cs.stanford.edu/~plogfren/bidirectional\\_ppr\\_thesis.pdf](https://cs.stanford.edu/~plogfren/bidirectional_ppr_thesis.pdf)

[3]. I. Manipur, et al., "Netpro2vec: A Graph Embedding Framework for Biomedical Applications" in IEEE/ACM Transactions on Computational Biology and Bioinformatics, vol. 19, no. 02

# What is Graph Random Walk (RW)

## Input


- A graph  $G$
- A set of walkers  $Q$ , with start vertices

## Walking Process

- Each walker selects a neighbor of current vertex at random
- Move to the selected neighbor
- Repeat until the termination condition is met

## Output

- Walking path of walkers in  $Q$



# What is Graph Random Walk (RW)

## Input


- A graph  $G$
- A set of walkers  $Q$ , with start vertices

## Walking Process

- Each walker selects a neighbor of current vertex at random
- Move to the selected neighbor
- Repeat until the termination condition is met

## Output

- Walking path of walkers in  $Q$



# What is Graph Random Walk (RW)

## Input


- A graph  $G$
- A set of walkers  $Q$ , with start vertices

## Walking Process

- Each walker selects a neighbor of current vertex at random
- Move to the selected neighbor
- Repeat until the termination condition is met

## Output

- Walking path of walkers in  $Q$



# Dynamic Graph Random Walk (DGRW)

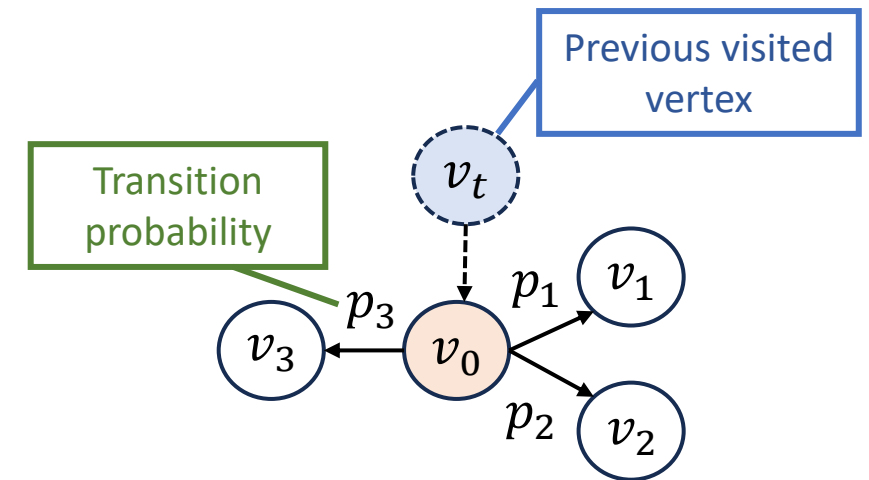
RW applications differ in the way of selecting neighbors, i.e., **how to calculate transition probability ( $p$ )**

## Static Graph Random Walk

- Calculate and store  $p$  **before random walk starts**
- $p$  is **fixed** during runtime

## Dynamic Graph Random Walk

- Calculate  $p$  dynamically **during runtime**
- $p$  might **change**
- Example: in Node2Vec<sup>[1]</sup>,  $\{p_1, p_2, p_3\}$  is dependent on the state of  $v_t$ , cannot be computed in advance



# How to Process DGRW

## DGRW

- Higher computation workload, but less space cost than SGRW
- A huge amount of queries



## CPU

- ✗ Limited computing cores

## FPGA

- ✗ Focus on hardware design

## GPU

- ✓ Massive computing cores
- ✓ Potential for high parallelism

# How to Process DGRW

## DGRW

- Higher computation workload
- Higher space cost than SGRW
- A huge amount of queries

### **Challenges**

- *Limited memory space*
- *Load imbalance*

## CPU

- ✗ Limited computing cores

## GPU

- ✓ Massive computing cores
- ✓ Potential for high parallelism

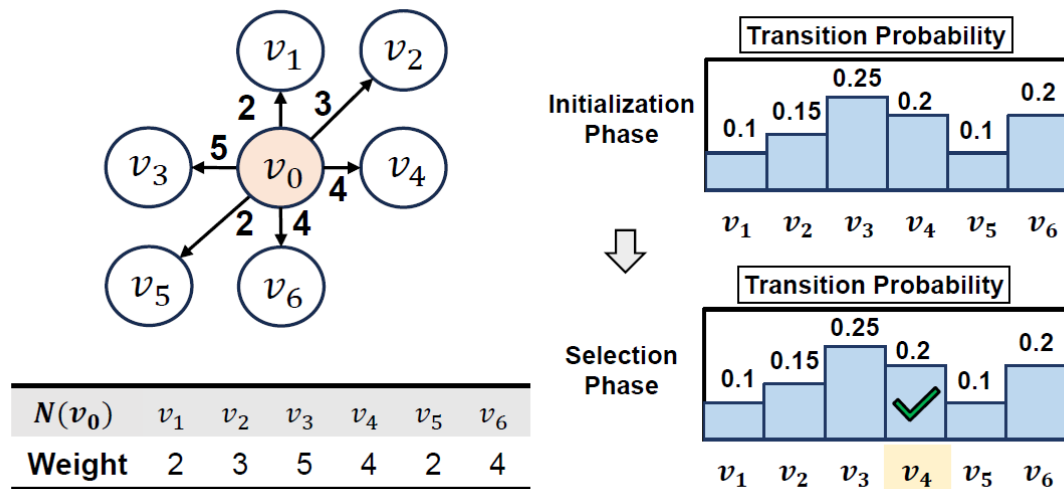
Hardware design



# Challenge: Limited Memory Space

## Memory cost of DGRW:

- **Graph data** – necessary data, cannot be eliminated
- **Walking trace** – necessary data, cannot be eliminated
- **Buffer for computation** – necessary data?
  - Sampling: initialization, then selection
  - Sampling algorithms, like ITS and ALS, require an  $O(d)$  buffer,  $d$  is the vertex degree




| Sampling method | ITS         | ALS    |
|-----------------|-------------|--------|
| Initialization  | $O(d)$      | $O(d)$ |
| Selection       | $O(\log d)$ | $O(1)$ |
| Space Cost      | $O(d)$      | $O(d)$ |

# Challenge: Limited Memory Space

## Memory cost of DGRW:

- **Graph data** – necessary data, cannot be eliminated
- **Walking trace** – necessary data, cannot be eliminated
- **Buffer for computation** – necessary data?
  - Sampling: initialization, then selection
  - Sampling algorithms, like ITS and ALS, requires an  $O(d)$  buffer,  $d$  is the vertex degree
  - Only a few queries can be hold on GPU simultaneously



**The Twitter Graph**  
( $d_{max} = 3 \times 10^6$ )

Memory overhead when  
 $N_{query} = 10^6$  in Skywalker:

- Graph data: **18GB**
- Trace: **0.3GB**
- A single query: **11.45MB**

*A few queries  
can be hold*



**A100**

Memory space of modern  
GPU is not very large:


- A100: **40GB/80GB**

# Challenge: Load Imbalance

## The sampling workload varies:

- Workload is governed by vertex degree
- Degrees in real-world graph follow **power-law distribution**
- Huge amounts of computing cores on GPU exacerbates the imbalance problem

**80%** of total vertices with lowest **28%** degrees



**14%** of total vertices with highest **20%** degrees

# Our Solution

## FlowWalker: an efficient DGRW framework at minimal memory cost

- Adopt **reservoir sampling** to reduce sampling space complexity to  $O(1)$ .



✓ Eliminate computation buffer

- High-performance processing engine, which leverages a **sampler-centric computation model** and performs **dynamic scheduling**.



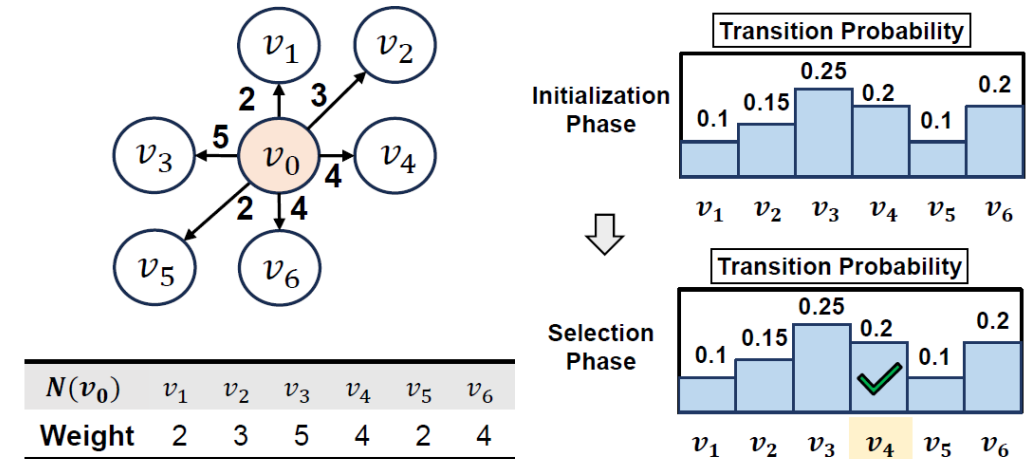
✓ Mitigate load imbalance

FlowWalker supports DGRW on large graphs with high parallelism.

# Reservoir Sampling (RS)

## RS saves memory:

- Sampling: initialization, then selection
- Sampling algorithms, like ITS and ALS, require an  $O(d)$  buffer,  $d$  is the vertex degree
- RS is pre-processing free and diminishes space complexity to  $O(1)$
- Can be accelerated through parallel processing



| Sampling method | ITS         | ALS    | RS     |
|-----------------|-------------|--------|--------|
| Initialization  | $O(d)$      | $O(d)$ | -      |
| Selection       | $O(\log d)$ | $O(1)$ | $O(d)$ |
| Space Cost      | $O(d)$      | $O(d)$ | $O(1)$ |

# FlowWalker: Parallelize Reservoir Sampling

## Direct Parallel Reservoir Sampling (DPRS)

**Input:** number of neighbors  $n$ ; number of threads  $k$ ;

**Output:** the selected vertex;

**for range**  $[0, \frac{n}{k} - 1]$ :

**parallel for**  $k$  neighbors of  $v$ : calculate bias once

$w_v = \text{GET\_EDGE\_BIAS}(v);$

$p_v = \text{PREFIX\_SUM}(w_v);$

*/\*Get the result of the local thread\*/* scan in every loop

$\text{select\_local} = \text{SELECTION}(p_v);$

*/\*Get the global result of this loop\*/* reduction in every loop

$\text{select\_global} = \text{REDUCTION}(\text{select\_local});$

*/\*Return the final result after all the loops\*/*

**return**  $\text{select\_global};$

- $O\left(\frac{n}{k} \times \log k\right)$  time,  $O(1)$  space
- Used when the bias calculation overhead is **high**  
(i.e. Node2Vec requires a binary search)

## Zig-Zag Parallel Reservoir Sampling (ZPRS)

**Input:** number of neighbors  $n$ ; number of threads  $k$ ;

**Output:** the selected vertex;

**for range**  $[0, \frac{n}{k} - 1]$ :

**parallel for**  $k$  neighbors of  $v$ :

$w_v += \text{GET\_EDGE\_BIAS}(v);$

$p_v = \text{PREFIX\_SUM}(w_v);$  one global scan

**for range**  $[0, \frac{n}{k} - 1]$ :

**parallel for**  $k$  neighbors of  $v$ :

$p_v += \text{GET\_EDGE\_BIAS}(v);$  calculate bias twice

$\text{select\_local} = \text{SELECTION}(p_v);$

$\text{select\_global} = \text{REDUCTION}(\text{select\_local});$  one global reduction

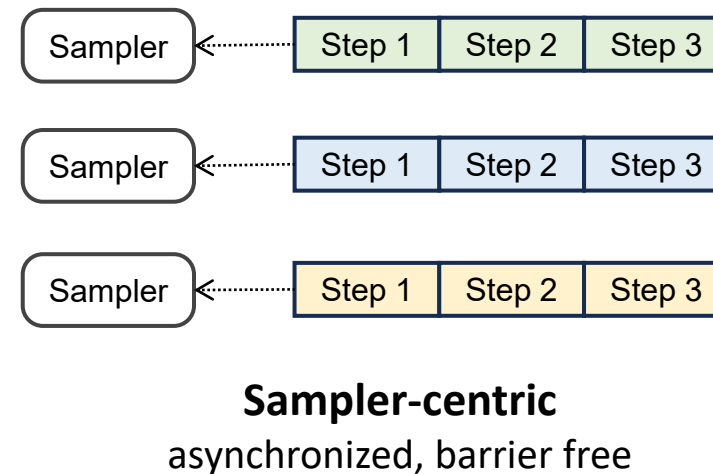

**return**  $\text{select\_global};$

- $O\left(\frac{n}{k} + \log k\right)$  time,  $O(1)$  space
- Used when the bias calculation overhead is **low**  
(i.e. MetaPath only needs a label matching)

# FlowWalker: Execution Engine

## Sampler-centric computation

- Step-centric model: **global barrier** between each step
- Sampler-centric model: **eliminate global barrier**
- Samplers with various computing capabilities cater for vertices with different degrees



# FlowWalker: Execution Engine

## Sampler-centric computation

- Step-centric model: **global barrier** between each step
- Sampler-centric model: **eliminate global barrier**
- Samplers with various computing capabilities cater for vertices with different degrees

## Dynamic scheduling

- Fetches one query from the global task pool when one query completes
- A walking query will not be evicted until it completes
- Achieves **adaptive load balancing**



# Experiment Setup

## Baseline

- **Skywalker** [*PACT' 21*] – GPU-based framework
- **LightRW** [*SIGMOD' 23*] – FPGA-based dynamic RW framework
- **ThunderRW** [*VLDB' 21*] – CPU in-memory framework
- **DGL** – widely adopted GNN framework, run in dynamic mode, CPU for Node2Vec, and GPU for other applications

## Datasets

- 10 read-world datasets, including **5 billion-scale datasets**

## Applications

- DeepWalk, Personalized PageRank (PPR), Node2Vec, MetaPath

## Environment

- A100 (40 GB) GPU, 100 KB shared memory of each SM
- AMD Alveo U250 FPGA
- Intel 8336C CPU with 16 cores and hyper-threading enabled

| Dataset     | Name | V     | E     | $d_{max}$ | Size(GB) |
|-------------|------|-------|-------|-----------|----------|
| com-youtube | YT   | 1.1 M | 6 M   | 28K       | 0.05     |
| cit-patents | CP   | 3.8 M | 33 M  | 793       | 0.26     |
| Livejournal | LJ   | 4.8 M | 86 M  | 20K       | 0.66     |
| Orkut       | OK   | 3.1 M | 234 M | 33K       | 1.76     |
| EU-2015     | EU   | 11 M  | 522M  | 399K      | 3.93     |
| Arabic-2005 | AB   | 23 M  | 1.1B  | 576K      | 8.34     |
| UK-2005     | UK   | 39 M  | 1.6B  | 1.7M      | 11.82    |
| Twitter     | TW   | 42 M  | 2.4 B | 3M        | 18.08    |
| Friendster  | FS   | 66 M  | 3.6 B | 5K        | 27.16    |
| SK-2005     | SK   | 51 M  | 3.6 B | 8.5M      | 27.16    |


## Datasets

# Overall Comparison

- **FlowWalker** is the only framework completing all test cases
- **Speedup**: FlowWalker achieves significant speedup, up-to 752.2x
- **Memory**: The extra memory cost of FlowWalker stays constant accross different datasets

| Framework | Platform | Maximum Speedup |
|-----------|----------|-----------------|
| DGL       | GPU      | 92.2x           |
| DGL       | CPU      | 315.8x          |
| LightRW   | FPGA     | 16.4x           |
| ThunderRW | CPU      | 752.2x          |
| Skywalker | GPU      | 72.1x           |


The speedup of FlowWalker over baselines



Extra memory cost ( $V_{total} - V_{dataset}$ ) of FlowWalker and Skywalker

# Case Study

- **Friend recommendation GNN** in Douyin (a popular video APP), implemented based on GraphLearn
- Test graph contains 227 million vertices and 2.71 billion edges
- FlowWalker reduces the RW time from **35%** (3.49 hours) to **3%** (13 minutes)



Training one epoch of GNN

# Conclusion

- FlowWalker is an efficient GPU dynamic graph random walk framework.
- FlowWalker employs the reservoir sampling and **reduce sampling memory cost to  $O(1)$** .
- FlowWalker uses dynamic walking engine and sampler-centric model to **mitigate workload imbalance and the global barrier**.
- FlowWalker achieves a significant speedup with minimal memory cost. It can process on large graphs with a high parallelism.



Scan and star 😊

Source code at <https://github.com/junyimei/flowwalker-artifact>

Contact: [meijunyi@sjtu.edu.cn](mailto:meijunyi@sjtu.edu.cn)

# Thank you!

---

## Q&A



Scan and star 😊

Source code at <https://github.com/junyimei/flowwalker-artifact>

Contact: [meijunyi@sjtu.edu.cn](mailto:meijunyi@sjtu.edu.cn)